

Project 1 Report

Alia Mohamed

March 19, 2023

Task 1: Launch a Kubernetes Cluster with one master and at least two workers.
 For this task I chose to work with AWS EKS and the command line tool KubeCTL.
 The following commands and output were used to setup and test KubeCTL:

```
kubectrl version --short --client
zsh: command not found: kubectrl
alia@Alias-MBP ~ % curl -O https://s3.us-west-2.amazonaws.com/amazon-eks/1.25.6/2023-01-30/bin/darwin/amd64/kubectrl
% Total % Received % Xferd Average Speed Time Time Time Current
Dload Upload Total Spent Left Speed
100 47.9M 100 47.9M 0 0 7729k 0 0:00:06 0:00:06 --:--:-- 8935k
alia@Alias-MBP ~ % curl -O https://s3.us-west-2.amazonaws.com/amazon-eks/1.25.6/2023-01-30/bin/darwin/amd64/kubectrl.sha256
% Total % Received % Xferd Average Speed Time Time Time Current
Dload Upload Total Spent Left Speed
100 73 100 73 0 0 203 0 --:--:-- --:--:-- --:--:-- 207
alia@Alias-MBP ~ % openssl sha1 -sha256 kubectrl
SHA256(kubectrl)= 9102988689ff79b44334c406d799ff2d1ec4afbe11ac5612c60b51f2e21be72f
alia@Alias-MBP ~ % chmod +x ./kubectrl
alia@Alias-MBP ~ % mkdir -p $HOME/bin && cp ./kubectrl $HOME/bin/kubectrl && export PATH=$HOME/bin:$PATH
alia@Alias-MBP ~ % echo 'export PATH=$PATH:$HOME/bin' >> ~/.bash_profile
alia@Alias-MBP ~ % kubectrl version --short --client
Flag --short has been deprecated, and will be removed in the future. The --short output will become the default.
Client Version: v1.25.6-eks-48e63af
Kustomize Version: v4.5.7
```

To launch the EKS cluster I used the command: `eksctl create cluster --name my-cluster --region us-east-1 --nodegroup-name my-worker-group --node-type t2.micro --nodes 2`

Output for the command:

```
2023-03-19 18:06:24 [i] eksctl version 0.134.0
2023-03-19 18:06:24 [i] using region us-east-1
2023-03-19 18:06:25 [i] setting availability zones to [us-east-1b us-east-1a]
2023-03-19 18:06:25 [i] subnets for us-east-1b - public:192.168.0.0/19 private:192.168.64.0/19
2023-03-19 18:06:25 [i] subnets for us-east-1a - public:192.168.32.0/19 private:192.168.96.0/19
2023-03-19 18:06:25 [i] nodegroup "my-worker-group" will use "" [AmazonLinux2/1.25]
2023-03-19 18:06:25 [i] using Kubernetes version 1.25
2023-03-19 18:06:25 [i] creating EKS cluster "my-cluster" in "us-east-1" region with managed nodes
2023-03-19 18:06:25 [i] will create 2 separate CloudFormation stacks for cluster itself and the initial managed nodegroup
2023-03-19 18:06:25 [i] if you encounter any issues, check CloudFormation console or try 'eksctl utils describe-stacks --region=us-east-1 --cluster=my-cluster'
2023-03-19 18:06:25 [i] Kubernetes API endpoint access will use default of {publicAccess=true, privateAccess=false} for cluster "my-cluster" in "us-east-1"
2023-03-19 18:06:25 [i] CloudWatch logging will not be enabled for cluster "my-cluster" in "us-east-1"
2023-03-19 18:06:25 [i] you can enable it with 'eksctl utils update-cluster-logging --enable-types={SPECIFY-YOUR-LOG-TYPES-HERE (e.g. all)} --region=us-east-1 --cluster=my-cluster'
2023-03-19 18:06:25 [i]
2 sequential tasks: { create cluster control plane "my-cluster",
  2 sequential sub-tasks: {
    wait for control plane to become ready,
    create managed nodegroup "my-worker-group",
  }
}
2023-03-19 18:06:25 [i] building cluster stack "eksctl-my-cluster-cluster"
2023-03-19 18:06:26 [i] deploying stack "eksctl-my-cluster-cluster"
2023-03-19 18:06:56 [i] waiting for CloudFormation stack "eksctl-my-cluster-cluster"
2023-03-19 18:07:26 [i] waiting for CloudFormation stack "eksctl-my-cluster-cluster"
```

```

2023-03-19 18:08:26 [i] waiting for CloudFormation stack "eksctl-my-cluster-cluster"
2023-03-19 18:09:26 [i] waiting for CloudFormation stack "eksctl-my-cluster-cluster"
2023-03-19 18:10:26 [i] waiting for CloudFormation stack "eksctl-my-cluster-cluster"
2023-03-19 18:11:27 [i] waiting for CloudFormation stack "eksctl-my-cluster-cluster"
2023-03-19 18:12:27 [i] waiting for CloudFormation stack "eksctl-my-cluster-cluster"
2023-03-19 18:13:27 [i] waiting for CloudFormation stack "eksctl-my-cluster-cluster"
2023-03-19 18:14:27 [i] waiting for CloudFormation stack "eksctl-my-cluster-cluster"
2023-03-19 18:15:27 [i] waiting for CloudFormation stack "eksctl-my-cluster-cluster"
2023-03-19 18:16:27 [i] waiting for CloudFormation stack "eksctl-my-cluster-cluster"
2023-03-19 18:17:28 [i] waiting for CloudFormation stack "eksctl-my-cluster-cluster"
2023-03-19 18:18:28 [i] waiting for CloudFormation stack "eksctl-my-cluster-cluster"
2023-03-19 18:20:30 [i] building managed nodegroup stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 18:20:30 [i] deploying stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 18:20:30 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 18:21:01 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 18:21:31 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 18:22:53 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 18:24:08 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 18:24:08 [i] waiting for the control plane to become ready
2023-03-19 18:24:08 [✓] saved kubeconfig as "/Users/alia/.kube/config"
2023-03-19 18:24:08 [i] no tasks
2023-03-19 18:24:08 [✓] all EKS cluster resources for "my-cluster" have been created
2023-03-19 18:24:08 [i] nodegroup "my-worker-group" has 2 node(s)
2023-03-19 18:24:08 [i] node "ip-192-168-12-77.ec2.internal" is ready
2023-03-19 18:24:08 [i] node "ip-192-168-40-75.ec2.internal" is ready
2023-03-19 18:24:08 [i] waiting for at least 2 node(s) to become ready in "my-worker-group"
2023-03-19 18:24:08 [i] nodegroup "my-worker-group" has 2 node(s)
2023-03-19 18:24:08 [i] node "ip-192-168-12-77.ec2.internal" is ready
2023-03-19 18:24:08 [i] node "ip-192-168-40-75.ec2.internal" is ready
2023-03-19 18:24:10 [i] kubectl command should work with "/Users/alia/.kube/config", try 'kubectl get nodes'
2023-03-19 18:24:10 [✓] EKS cluster "my-cluster" in "us-east-1" region is ready

```

This operation took 18 minutes to complete resulting the creation of the cluster with two worker nodes and active EC2 instances

The following commands and output were used to obtain information about the cluster and nodes:

Command: `kubectl get nodes -o wide`

Output:

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP	OS-
IMAGE	KERNEL-VERSION	CONTAINER-RUNTIME					
ip-192-168-12-77.ec2.internal	Ready	<none>	7m36s	v1.25.6-eks-48e63af	192.168.12.77	18.234.214.145	Amazon Linux
2 5.10.167-147.601.amzn2.x86_64	containerd://1.6.6						
ip-192-168-40-75.ec2.internal	Ready	<none>	7m39s	v1.25.6-eks-48e63af	192.168.40.75	3.88.115.10	Amazon Linux
2 5.10.167-147.601.amzn2.x86_64	containerd://1.6.6						

Command: `kubectl get pods -A -o wide`

Output:

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED
node	READINESS GATES							
kube-system	aws-node-2rt27	1/1	Running	0	8m12s	192.168.40.75	ip-192-168-40-75.ec2.internal	
	<none>	<none>						
kube-system	aws-node-7q7fk	1/1	Running	0	8m9s	192.168.12.77	ip-192-168-12-77.ec2.internal	
	<none>	<none>						
kube-system	coredns-7975d6fb9b-8dvrb	1/1	Running	0	16m	192.168.59.169	ip-192-168-40-75.ec2.internal	
	<none>	<none>						
kube-system	coredns-7975d6fb9b-shv2b	1/1	Running	0	16m	192.168.60.116	ip-192-168-40-75.ec2.internal	
	<none>	<none>						
kube-system	kube-proxy-cvg8j	1/1	Running	0	8m12s	192.168.40.75	ip-192-168-40-75.ec2.internal	
	<none>	<none>						

```
kube-system kube-proxy-dgssd 1/1 Running 0 8m9s 192.168.12.77 ip-192-168-12-77.ec2.internal <none> <none>
```

Task 2: Deploy a containerized application which run multiple instances of the same container
For this task there were multiple approaches. The first approach was following the instruction on the EKS documentation as follows:

Creating a namespace to allow the creation of the application and group resources in the cluster:

Command: `kubectl create namespace my-kubernetes-namespace`

Output:

`namespace/my-kubernetes-namespace created`

creating deployment manifest using the code: this is the code for the deployment file for the application.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: eks-sample-linux-deployment
  namespace: my-kubernetes-namespace
  labels:
    app: eks-sample-linux-app
spec:
  replicas: 3
  selector:
    matchLabels:
      app: eks-sample-linux-app
  template:
    metadata:
      labels:
        app: eks-sample-linux-app
    spec:
      affinity:
        nodeAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            nodeSelectorTerms:
              - matchExpressions:
                  - key: kubernetes.io/arch
                    operator: In
                    values:
                      - amd64
                      - arm64
      containers:
        - name: nginx
          image: public.ecr.aws/nginx/nginx:1.21
          ports:
            - name: http
```

```

    containerPort: 80
    imagePullPolicy: IfNotPresent
  nodeSelector:
    kubernetes.io/os: linux

```

Applying the deployment manifest to the cluster:

Command: `kubectl apply -f eks-sample-deployment.yaml`

Output:

deployment.apps/eks-sample-linux-deployment created

Creating a service with the code: this is the code for the service provided by the application.

```

apiVersion: v1
kind: Service
metadata:
  name: eks-sample-linux-service
  namespace: my-kubernetes-namespace
  labels:
    app: eks-sample-linux-app
spec:
  selector:
    app: eks-sample-linux-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80

```

applying the server manifest to the cluster:

Command: `kubectl apply -f eks-sample-service.yaml`

Output:

service/eks-sample-linux-service created

Viewing resources of the namespace:

Command: `kubectl get all -n my-kubernetes-namespace`

Output:

```

NAME                                READY STATUS    RESTARTS AGE
pod/eks-sample-linux-deployment-7f646d456c-84jfx 1/1   Running      0      33s
pod/eks-sample-linux-deployment-7f646d456c-ktg2p 0/1   Pending      0      33s
pod/eks-sample-linux-deployment-7f646d456c-nbkw7 0/1   ContainerCreating 0      33s
NAME                                TYPE          CLUSTER-IP    EXTERNAL-IP  PORT(S) AGE
service/eks-sample-linux-service    ClusterIP     10.100.241.255 <none>       80/TCP    14s
NAME                                READY UP-TO-DATE AVAILABLE AGE
deployment.apps/eks-sample-linux-deployment 1/3   3          1      33s
NAME                                DESIRED CURRENT READY AGE
replicaset.apps/eks-sample-linux-deployment-7f646d456c 3      3          1      33s

```

Details of the service:

Command: `kubectl -n my-kubernetes-namespace describe service eks-sample-linux-service`

Output:

```

Name:      eks-sample-linux-service
Namespace:  my-kubernetes-namespace
Labels:    app=eks-sample-linux-app
Annotations: <none>
Selector:  app=eks-sample-linux-app
Type:      ClusterIP
IP Family Policy: SingleStack
IP Families: IPv4
IP:        10.100.241.255
IPs:       10.100.241.255
Port:      <unset> 80/TCP
TargetPort: 80/TCP
Endpoints: 192.168.23.117:80,192.168.4.191:80
Session Affinity: None
Events:    <none>

```

Details for one of the pods:

Command: kubectl -n my-kubernetes-namespace describe pod eks-sample-linux-deployment-7f646d456c-84jfx

Output:

```

Name:      eks-sample-linux-deployment-7f646d456c-84jfx
Namespace:  my-kubernetes-namespace
Priority:    0
Service Account: default
Node:      ip-192-168-12-77.ec2.internal/192.168.12.77
Start Time: Sun, 19 Mar 2023 18:31:40 -0400
Labels:    app=eks-sample-linux-app
           pod-template-hash=7f646d456c
Annotations: <none>
Status:     Running
IP:         192.168.4.191
IPs:        192.168.4.191
Controlled By: ReplicaSet/eks-sample-linux-deployment-7f646d456c
Containers:
  nginx:
    Container ID: containerd://e6f861822c3854de7b768987875a7cfd6a646e5b902b3f99e73590a8812df6eb
    Image:      public.ecr.aws/nginx/nginx:1.21
    Image ID:   public.ecr.aws/nginx/nginx@sha256:3aac7c736093ce043a17d6e83ef5addb8be321b5b6b93879141e51474448ca65
    Port:      80/TCP
    Host Port: 0/TCP
    State:     Running
      Started: Sun, 19 Mar 2023 18:31:46 -0400
    Ready:     True
    Restart Count: 0
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-29zkz (ro)
Conditions:
  Type           Status
  Initialized     True
  Ready          True
  ContainersReady True
  PodScheduled   True
Volumes:
  kube-api-access-29zkz:
    Type:      Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI: true
  QoS Class:      BestEffort
  Node-Selectors:  kubernetes.io/os=linux
  Tolerations:    node.kubernetes.io/not-ready:NoExecute op=Exists for 300s

```

node.kubernetes.io/unreachable:NoExecute op=Exists for 300s

Events:

Type	Reason	Age	From	Message
Normal	Scheduled	2m19s	default-scheduler	Successfully assigned my-kubernetes-namespace/eks-sample-linux-deployment-7f646d456c-84jfx to ip-192-168-12-77.ec2.internal
Normal	Pulling	2m18s	kubelet	Pulling image "public.ecr.aws/nginx/nginx:1.21"
Normal	Pulled	2m14s	kubelet	Successfully pulled image "public.ecr.aws/nginx/nginx:1.21" in 4.265129131s (4.265166386s including waiting)
Normal	Created	2m14s	kubelet	Created container nginx
Normal	Started	2m13s	kubelet	Started container nginx

The second approach was to use another service application hamster.yaml

Deploying the application:

Command: kubectl apply -f examples/hamster.yaml

Output:

verticalpodautoscaler.autoscaling.k8s.io/hamster-vpa created
deployment.apps/hamster created

viewing pods for the application

Command: kubectl get pods -l app=hamster

Output:

NAME	READY	STATUS	RESTARTS	AGE
hamster-59cc68d575-hhpsj	0/1	Pending	0	27s
hamster-59cc68d575-xd9wf	0/1	Pending	0	27s

description for one of the pods:

Command: kubectl describe pod hamster-59cc68d575-hhpsj

Output:

```
Name:          hamster-59cc68d575-hhpsj
Namespace:    default
Priority:      0
Service Account: default
Node:         <none>
Labels:       app=hamster
              pod-template-hash=59cc68d575
Annotations:  <none>
Status:       Pending
IP:
IPs:          <none>
Controlled By: ReplicaSet/hamster-59cc68d575
Containers:
  hamster:
    Image:      registry.k8s.io/ubuntu-slim:0.1
    Port:       <none>
    Host Port:  <none>
    Command:
      /bin/sh
    Args:
      -c
      while true; do timeout 0.5s yes >/dev/null; sleep 0.5s; done
    Requests:
      cpu:      100m
      memory:   50Mi
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-66r2g (ro)
Conditions:
  Type          Status
  PodScheduled  False
```

Volumes:
 kube-api-access-66r2g:
 Type: Projected (a volume that contains injected data from multiple sources)
 TokenExpirationSeconds: 3607
 ConfigMapName: kube-root-ca.crt
 ConfigMapOptional: <nil>
 DownwardAPI: true
 QoS Class: Burstable
 Node-Selectors: <none>
 Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
 node.kubernetes.io/unreachable:NoExecute op=Exists for 300s

Events:

Type	Reason	Age	From	Message
Warning	FailedScheduling	67s	default-scheduler	0/2 nodes are available: 2 Too many pods. preemption: 0/2 nodes are available: 2 No preemption victims found for incoming pod.

All operation during this task did not take a very long time to execute. All operation almost executed instantly.

Task 3: Scale the pod to more container instances.

Looking through the documentation provided for EKS, the options that I found for scaling the cluster were the Vertical and Horizontal Autoscalers which I used for this part of the project.

Deploying the Vertical autoscaler: this step will allow the cluster to grow vertically creating more space and resources for the application

Command: vertical-pod-autoscaler % ./hack/vpa-up.sh

Output:

```
customresourcedefinition.apiextensions.k8s.io/verticalpodautoscalercheckpoints.autoscaling.k8s.io created
customresourcedefinition.apiextensions.k8s.io/verticalpodautoscalers.autoscaling.k8s.io created
clusterrole.rbac.authorization.k8s.io/system:metrics-reader created
clusterrole.rbac.authorization.k8s.io/system:vpa-actor created
clusterrole.rbac.authorization.k8s.io/system:vpa-checkpoint-actor created
clusterrole.rbac.authorization.k8s.io/system:evictioner created
clusterrolebinding.rbac.authorization.k8s.io/system:metrics-reader created
clusterrolebinding.rbac.authorization.k8s.io/system:vpa-actor created
clusterrolebinding.rbac.authorization.k8s.io/system:vpa-checkpoint-actor created
clusterrole.rbac.authorization.k8s.io/system:vpa-target-reader created
clusterrolebinding.rbac.authorization.k8s.io/system:vpa-target-reader-binding created
clusterrolebinding.rbac.authorization.k8s.io/system:vpa-evictioner-binding created
serviceaccount/vpa-admission-controller created
serviceaccount/vpa-recommender created
serviceaccount/vpa-updater created
clusterrole.rbac.authorization.k8s.io/system:vpa-admission-controller created
clusterrolebinding.rbac.authorization.k8s.io/system:vpa-admission-controller created
clusterrole.rbac.authorization.k8s.io/system:vpa-status-reader created
clusterrolebinding.rbac.authorization.k8s.io/system:vpa-status-reader-binding created
deployment.apps/vpa-updater created
deployment.apps/vpa-recommender created
Generating certs for the VPA Admission Controller in /tmp/vpa-certs.
Generating RSA private key, 2048 bit long modulus
.....+++
.....+++
e is 65537 (0x10001)
unknown option -addext
req [options] <infile >outfile
where options are
-inform arg  input format - DER or PEM
-outform arg output format - DER or PEM
-in arg     input file
-out arg    output file
```

```

-text      text form of request
-pubkey    output public key
-noout     do not output REQ
-verify    verify signature on REQ
-modulus   RSA modulus
-nodes     don't encrypt the output key
-subject   output the request's subject
-passin    private key password source
-key file  use the private key contained in file
-keyform arg key file format
-keyout arg file to send the key to
-newkey rsa:bits generate a new RSA key of 'bits' in size
-newkey dsa:file generate a new DSA key, parameters taken from CA in 'file'
-newkey ec:file generate a new EC key, parameters taken from CA in 'file'
-[digest]  Digest to sign with (md5, sha1, md4)
-config file request template file.
-subj arg  set or modify request subject
-multivalue-rdn enable support for multivalued RDNs
-new       new request.
-batch     do not ask anything during request generation
-x509      output a x509 structure instead of a cert. req.
-days      number of days a certificate generated by -x509 is valid for.
-set serial serial number to use for a certificate generated by -x509.
-newhdr     output "NEW" in the header lines
-asn1-kludge Output the 'request' in a format that is wrong but some CA's
            have been reported as requiring
-extensions .. specify certificate extension section (override value in config file)
-reqexts ..  specify request extension section (override value in config file)
-utf8      input characters are UTF8 (default ASCII)
-nameopt arg - various certificate name options
-reqopt arg - various request text options

```

ERROR: Failed to create CA certificate for self-signing. If the error is "unknown option -addext", update your openssl version or deploy VPA from the vpa-release-0.8 branch.
 deployment.apps/vpa-admission-controller created
 service/vpa-webhook created

Verifying that the autoscaler works:

Command: `kubectl get pods -n kube-system`

Output:

NAME	READY	STATUS	RESTARTS	AGE
aws-node-2rt27	1/1	Running	0	120m
aws-node-7q7fk	1/1	Running	0	120m
coredns-7975d6fb9b-8dvrb	1/1	Running	0	128m
coredns-7975d6fb9b-shv2b	1/1	Running	0	128m
kube-proxy-cvg8j	1/1	Running	0	120m
kube-proxy-dgssd	1/1	Running	0	120m
metrics-server-8ff8f88c6-4zg99	1/1	Running	0	97m
vpa-admission-controller-5cc669f7f8-6v28m	0/1	Pending	0	2m19s
vpa-recommender-5bcd4fd7bc-gzgpx	0/1	Pending	0	2m20s
vpa-updater-68c46997f8-278vx				

Describing the VPA for the scaler:

Command: `kubectl describe vpa/hamster-vpa`

Output:

```

Name:      hamster-vpa
Namespace: default
Labels:    <none>
Annotations: <none>
API Version: autoscaling.k8s.io/v1
Kind:      VerticalPodAutoscaler
Metadata:
  Creation Timestamp: 2023-03-20T00:23:29Z
  Generation:        1

```



```

Managed Fields:
API Version: autoscaling.k8s.io/v1
Fields Type: FieldsV1
fieldsV1:
f:metadata:
  f:annotations:
    .:
    f:kubectl.kubernetes.io/last-applied-configuration:
f:spec:
  .:
  f:resourcePolicy:
    .:
    f:containerPolicies:
  f:targetRef:
Manager:      kubectl-client-side-apply
Operation:    Update
Time:         2023-03-20T00:23:29Z
Resource Version: 28782
UID:          98ff8a63-d6cd-402c-b34d-56bac2ff6260
Spec:
Resource Policy:
Container Policies:
  Container Name: *
  Controlled Resources:
    cpu
    memory
Max Allowed:
  Cpu: 1
  Memory: 500Mi
Min Allowed:
  Cpu: 100m
  Memory: 50Mi
Target Ref:
  API Version: apps/v1
  Kind: Deployment
  Name: hamster
Events:      <none>

```

Creating a horizontal autoscaler: this step will allow the cluster to grow horizontally providing more resources for the services.

Command: `kubectl autoscale deployment hamster --cpu-percent=50 --min=1 --max=10`

Output:

horizontalpodautoscaler.autoscaling/hamster autoscaled

Describing the autoscaler:

Command: `kubectl get hpa`

Output:

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
hamster	Deployment/hamster	<unknown>/50%	1	10	2	61s

for this task I noticed that the autoscores work to expand the network as needed when a service that need more CPU and Memory is deployed by creating more replicas and instances.

Task 4: Update the application with a new software version

For this task I chose to update the hamster.yaml application which was already up to date when it was first installed.

Command: `kubectl patch deployment hamster -p '{"space": {"minReadySeconds": 20}}'`

Output:

Warning: unknown field "space"

deployment.apps/hamster patched (no change)

Task 5: delete the application and stop the cluster

Deleting the application: deleting the application file that were created.

Command: `kubectl delete -f examples/hamster.yaml`

Output:

verticalpodautoscaler.autoscaling.k8s.io "hamster-vpa" deleted

deployment.apps "hamster" deleted

Stopping the cluster: deleting the cluster and all the resources and file that were created.

Command: `delete cluster --name my-cluster`

Output:

```
2023-03-19 21:01:56 [i] deleting EKS cluster "my-cluster"
2023-03-19 21:01:56 [i] will drain 0 unmanaged nodegroup(s) in cluster "my-cluster"
2023-03-19 21:01:56 [i] starting parallel draining, max in-flight of 1
2023-03-19 21:01:57 [i] deleting Fargate profile "my-fargate-profile"
2023-03-19 21:06:14 [i] deleted Fargate profile "my-fargate-profile"
2023-03-19 21:06:14 [i] deleted 1 Fargate profile(s)
2023-03-19 21:06:14 [i] will delete stack "eksctl-my-cluster-fargate"
2023-03-19 21:06:14 [✓] kubeconfig has been updated
2023-03-19 21:06:14 [i] cleaning up AWS load balancers created by Kubernetes objects of Kind Service or Ingress
2023-03-19 21:06:16 [i]
2 sequential tasks: { delete nodegroup "my-worker-group", delete cluster control plane "my-cluster" [async]
}
2023-03-19 21:06:16 [i] will delete stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 21:06:16 [i] waiting for stack "eksctl-my-cluster-nodegroup-my-worker-group" to get deleted
2023-03-19 21:06:16 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 21:06:46 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 21:07:30 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 21:08:52 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 21:09:53 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 21:11:07 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 21:12:26 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 21:14:05 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 21:15:53 [i] waiting for CloudFormation stack "eksctl-my-cluster-nodegroup-my-worker-group"
2023-03-19 21:15:53 [i] will delete stack "eksctl-my-cluster-cluster"
2023-03-19 21:15:54 [✓] all cluster resources were deleted
```

References

Installing or updating kubectl from <https://docs.aws.amazon.com/eks/latest/userguide/install-kubectl.html>

Getting started with Amazon EKS – eksctl from <https://docs.aws.amazon.com/eks/latest/userguide/getting-started-eksctl.html>

AWS Fargate Profile from <https://docs.aws.amazon.com/eks/latest/userguide/fargate-profile.html>

Deploy a sample application from <https://docs.aws.amazon.com/eks/latest/userguide/sample-deployment.html>

Installing the Kubernetes Metrics Server from <https://docs.aws.amazon.com/eks/latest/userguide/metrics-server.html>

Horizontal Pod Autoscaler from <https://docs.aws.amazon.com/eks/latest/userguide/horizontal-pod-autoscaler.html>

Vertical Pod Autoscaler from <https://docs.aws.amazon.com/eks/latest/userguide/vertical-pod-autoscaler.html>